

Project 3 - Machine Learning - Hand Gestures

- [Enunciado \(ENUNCIADO.pdf\)](#)

Instructions

The project instructions are located at [ENUNCIADO.pdf \(ENUNCIADO.pdf\)](#).

Installation

Instructions for installing OpenCV on Ubuntu 14.04 or 13.10: [here \(http://www.sysads.co.uk/2014/05/install-opencv-2-4-9-ubuntu-14-04-13-10/\)](http://www.sysads.co.uk/2014/05/install-opencv-2-4-9-ubuntu-14-04-13-10/).

Instructions for installing OpenCV on Ubuntu 13.04 or below: [here \(https://help.ubuntu.com/community/OpenCV\)](https://help.ubuntu.com/community/OpenCV).

Download the [example sets \(ftp://mi.eng.cam.ac.uk/pub/CamGesData/\)](ftp://mi.eng.cam.ac.uk/pub/CamGesData/) and extract them to the examples folder.

Generating ARFF files

Run `arff_generator.py`, providing the number of frames extracted per example (more than 2 -- more frames provides better but slower classification) and the example sets that will be used.

The following call extracts 6 frames per example, uses the Sets 2, 3 and 5, and saves it all on a file:

```
$ python extractor/arff_generator.py 6 2 3 5 > 6-235.arff
```

Features

For each example, we extract a number of frames from it.

For example, for 2 frames, we get the first and last frames.

For 3 frames, we get the first, the last and the middle frames.

Then, we extract features for each of these individual frames.

The frame image is segmented based on a HSV range roughly corresponding to skin color range: [0, 30, 60] to [50, 150, 255];

Then, from all the contours in the image (found with `findContours`), we grab the one with the biggest area.

We use this contour to fit a minimum-area rectangle, an ellipse, and a convex hull. We also extract convexity defects between the contour and the convex hull, and the contour's Hu moments.

From these primitives, we can extract all of the features:

```
area
perimeter
convex_hull_area
solidity
rect_center_x
rect_center_y
rect_width
rect_height
rect_angle
rect_aspect_ratio
ellipse_center_x
ellipse_center_y
ellipse_major_axis
ellipse_minor_axis
ellipse_angle
farthest_convex_defect
hu_moment_1
hu_moment_2
```

```
hu_moment_3
hu_moment_4
hu_moment_5
hu_moment_6
hu_moment_7
```

These 23 features are extracted for each frame. So, if an example is composed of 10 frames, it'll actually have 230 features.

In the ARFF file, we append a number to each attribute, corresponding to the frame number.

Functions

Based on the file generated by the following run, containing 2 frames per example and all 5 sets:

```
$ python extractor/arff_generator.py 2 1 2 3 4 5 > All.arff
```

We tested some classification functions and varied their parameters.

Using all 5 sets is the most challenging, because of the illumination setup by each set.

Extracting 2 frames per example improves the runtime of the classifier. Had we used more frames, we'd have a better result overall. Because right now our objective is to find the best function, we can use only 2.

Decision Tree (trees/J48)

Parameters

binarySplits had no influence on the results.

confidenceFactor did not impact much on classification performance, but higher levels slowed down the classification quite a bit. We maintained it at 0.25.

minNumObj was optimal at 10.

reducedErrorPruning deemed worse results, so we turned it off.

Results

Time taken to build model: 0.33 seconds		
Correctly Classified Instances	634	70.4444 %
Incorrectly Classified Instances	266	29.5556 %
Kappa statistic	0.6675	
Mean absolute error	0.0876	
Root mean squared error	0.2293	
Relative absolute error	44.3687 %	
Root relative squared error	72.971 %	
Total Number of Instances	900	

K-Nearest Neighbor (lazy/IBk)

Parameters

KNN was optimal at 1.

distanceWeighting had no influence on the results.

meanSquares had no influence on the results.

nearestNeighbourSearchAlgorithm we tried different ones, but all yielded the same result.

windowSize was only good at 0.

Results

Time taken to build model: 0 seconds		
--------------------------------------	--	--

Correctly Classified Instances	712	79.1111 %
Incorrectly Classified Instances	188	20.8889 %
Kappa statistic	0.765	
Mean absolute error	0.0481	
Root mean squared error	0.2143	
Relative absolute error	24.3407 %	
Root relative squared error	68.1877 %	
Total Number of Instances	900	

SVM (functions/libsvm)

Parameters

SVMType nu-CSV was better than C-CSV.

coef0 1.0 (small influence).

cost 1.0 (no influence).

degrees was optimal at 3

nu was optimal at 0.1

Results

Time taken to build model: 22.3 seconds		
Correctly Classified Instances	771	85.6667 %
Incorrectly Classified Instances	129	14.3333 %
Kappa statistic	0.8388	
Mean absolute error	0.0363	
Root mean squared error	0.1603	
Relative absolute error	18.3595 %	
Root relative squared error	51.0225 %	
Total Number of Instances	900	

Neural Network (functions/MultilayerPerceptron)

Parameters

hiddenLayers t (attrs+classes)

learningRate 0.3

momentum 0.5

trainingTime 500 (diminishing returns).

Results

Time taken to build model: 1.75 seconds		
Correctly Classified Instances	780	86.6667 %
Incorrectly Classified Instances	120	13.3333 %
Kappa statistic	0.85	
Mean absolute error	0.0296	
Root mean squared error	0.1721	
Relative absolute error	15 %	
Root relative squared error	54.7723 %	
Total Number of Instances	900	

Maximizing

Now that we know that, for this case, using a NN works best, we're going to try to achieve the highest possible correct classification rate.

For this, we'll use only the 1st set, extracting 10 frames from each example.

Results

Time taken to build model: 89.95 seconds		
Correctly Classified Instances	173	96.1111 %
Incorrectly Classified Instances	7	3.8889 %
Kappa statistic	0.9563	
Mean absolute error	0.014	
Root mean squared error	0.0845	
Relative absolute error	7.074 %	
Root relative squared error	26.8859 %	
Total Number of Instances	180	

Implementing

We implemented a Neural Network algorithm with backpropagation in Ruby.

It can be used like this:

```
require './nn'

xor = [
  [[0,0], [0]],
  [[0,1], [1]],
  [[1,0], [1]],
  [[1,1], [0]]
]

# input, hidden, output
n = NeuralNetwork.new(2, 2, 1)

puts "Training..."
n.train(xor)

puts "Evaluating..."
correct = n.test(xor)

puts "Got #{correct}/#{xor.length} -> #{100.0*correct/xor.length}%"
```

There's also a script that can take an ARFF file and pass it through the Neural Network.

```
$ ruby arff_nn.rb ../examples/Set1-2frames.arff
Training...
Evaluating...
Got 162/180 -> 90%
```

(This takes a long time! Our implementation is not that sophisticated)